

Cluster_Analysis__using_Kmeans

July 7, 2020

0.1 Introduction

K-means clustering algorithm represents each cluster by its corresponding cluster centroid. The algorithm will partition the input data data into k-disjoint clusters by using the following steps:

1. Form K clusters by assigning each instance to its nearest centroid.
2. Recompute the centroid of each cluster.

```
[1]: import pandas as pd
      %config IPCompleter.greedy=True
```

```
[2]: ratings=[[ 'John',5,5,2,1],
               ↪ ['Mary',4,5,3,2], ['Bob',4,4,4,3], ['Lisa',2,2,4,5], ['Lee',1,2,3,4], ['Harry',2,1,5,5]]
```

```
[3]: ratings
```

```
[3]: [['John', 5, 5, 2, 1],
      ['Mary', 4, 5, 3, 2],
      ['Bob', 4, 4, 4, 3],
      ['Lisa', 2, 2, 4, 5],
      ['Lee', 1, 2, 3, 4],
      ['Harry', 2, 1, 5, 5]]
```

```
[4]: titles=[ 'User', 'Jaws', 'Star Wars', 'Exorcist', 'Omen']
```

```
[5]: titles
```

```
[5]: ['User', 'Jaws', 'Star Wars', 'Exorcist', 'Omen']
```

```
[6]: movies=pd.DataFrame(ratings,columns=titles)
```

```
[7]: movies
```

```
[7]:
```

	User	Jaws	Star Wars	Exorcist	Omen
0	John	5	5	2	1
1	Mary	4	5	3	2
2	Bob	4	4	4	3
3	Lisa	2	2	4	5
4	Lee	1	2	3	4

5 Harry 2 1 5 5

0.2 Data Discusion

From the dataset the first 3 users like action movies, and the last 3 users enjoy horror movies. Our goal is to apply K-means clustering on the users to identify groups of users with similar movie preferences. K=2 from the dataset.

```
[8]: from sklearn import cluster
```

```
[9]: data=movies.drop('User',axis=1)
```

```
[10]: data
```

```
[10]:
```

	Jaws	Star Wars	Exorcist	Omen
0	5	5	2	1
1	4	5	3	2
2	4	4	4	3
3	2	2	4	5
4	1	2	3	4
5	2	1	5	5

```
[11]: k_means=cluster.KMeans(n_clusters=2,max_iter=50,random_state=1)
```

```
[12]: k_means
```

```
[12]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=50,
            n_clusters=2, n_init=10, n_jobs=None, precompute_distances='auto',
            random_state=1, tol=0.0001, verbose=0)
```

```
[13]: k_means.fit(data)
```

```
[13]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=50,
            n_clusters=2, n_init=10, n_jobs=None, precompute_distances='auto',
            random_state=1, tol=0.0001, verbose=0)
```

```
[14]: labels = k_means.labels_
```

```
[15]: labels
```

```
[15]: array([0, 0, 0, 1, 1, 1])
```

```
[16]: pd.DataFrame(labels,index=movies.User,columns=['Cluster ID'])
```

```
[16]:
```

	Cluster ID
User	
John	0

Mary	0
Bob	0
Lisa	1
Lee	1
Harry	1

K-means clustering has assigned the first 3 users to one cluster and the last 3 users to the second cluster. These results are consistent with our expectations. We can also display the centroid for each of the two clusters.

```
[17]: centroids=k_means.cluster_centers_
```

```
[18]: centroids
```

```
[18]: array([[4.33333333, 4.66666667, 3.          , 2.          ],
            [1.66666667, 1.66666667, 4.          , 4.66666667]])
```

```
[19]: pd.DataFrame(centroids,columns=data.columns)
```

```
[19]:
```

	Jaws	Star Wars	Exorcist	Omen
0	4.333333	4.666667	3.0	2.000000
1	1.666667	1.666667	4.0	4.666667

The cluster centroids can be used to determine other users cluster assignments.

```
[20]: import numpy as np
```

```
[21]: testData = np.array([[4,5,1,2],[3,2,4,4],[2,3,4,1],[3,2,3,3],[5,4,1,4]])
```

```
[22]: testData
```

```
[22]: array([[4, 5, 1, 2],
            [3, 2, 4, 4],
            [2, 3, 4, 1],
            [3, 2, 3, 3],
            [5, 4, 1, 4]])
```

```
[23]: labels=k_means.predict(testData)
```

```
[24]: labels
```

```
[24]: array([0, 1, 0, 1, 0])
```

```
[25]: labels=labels.reshape(-1,1)
```

```
[26]: labels
```

```
[26]: array([[0],
          [1],
          [0],
          [1],
          [0]])
```

```
[27]: usernames=np.array(['Paul','Kim','Liz','Tom','Bill']).reshape(-1,1)
```

```
[28]: cols=movies.columns.tolist()
```

```
[29]: cols
```

```
[29]: ['User', 'Jaws', 'Star Wars', 'Exorcist', 'Omen']
```

```
[30]: cols.append('Cluster ID')
```

```
[31]: cols
```

```
[31]: ['User', 'Jaws', 'Star Wars', 'Exorcist', 'Omen', 'Cluster ID']
```

```
[32]: newusers=pd.DataFrame(np.concatenate((usernames,testData,labels),
↪axis=1),columns=cols)
```

```
[33]: newusers
```

```
[33]:
```

	User	Jaws	Star Wars	Exorcist	Omen	Cluster ID
0	Paul	4	5	1	2	0
1	Kim	3	2	4	4	1
2	Liz	2	3	4	1	0
3	Tom	3	2	3	3	1
4	Bill	5	4	1	4	0

```
[34]: data
```

```
[34]:
```

	Jaws	Star Wars	Exorcist	Omen
0	5	5	2	1
1	4	5	3	2
2	4	4	4	3
3	2	2	4	5
4	1	2	3	4
5	2	1	5	5

```
[40]: centroids=k_means.cluster_centers_
```


0.3 Future Work

How to determine the number of clusters. We looked at the data and concluded that 2 clusters are okay. How about if you have a lot of data, or high dimensional data

```
[39]: import matplotlib.pyplot as plt
```

```
[41]: plt.scatter(centroids[:,0],centroids[:,1],c='black', s=200,alpha=.5)
```

```
[41]: <matplotlib.collections.PathCollection at 0x1e519cb3708>
```



```
[43]: k_means.inertia_
```

```
[43]: 9.333333333333334
```

1 How To Determine The Number of *Clusters*

We can determine the number of clusters in the data using the k-means clustering by varying the number of clusters within a range (this is a trial and error method). For example in this case we can vary from 1 to 6, and then compute the sum-of-squared-errors(SSE). The elbow in the plot of the SSE versus the number of clusters can be used to estimate the number of clusters.

```
[45]: import matplotlib.pyplot as plt
      %matplotlib inline
```

```
[46]: numClusters=[1,2,3,4,5,6]
```

```
[47]: SSE=[]
```

```
[51]: data
```

```
[51]:
```

	Jaws	Star Wars	Exorcist	Omen
0	5	5	2	1
1	4	5	3	2
2	4	4	4	3
3	2	2	4	5
4	1	2	3	4
5	2	1	5	5

```
[50]: for k in numClusters:
      k_means=cluster.KMeans(n_clusters=k)
      k_means.fit(data)
      SSE.append(k_means.inertia_)
```

```
[ ]:
```